

# Problem Set #3 (BDAT 1004)

By Daryoush Shabahang

August 19, 2021

## Question 1 Occupations - Python with Pandas

```
In [161]: # Step 1. Import the necessary libraries
import pandas as pd

# Step 2. Import the dataset from this address.
df1 = pd.read_csv(r'https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.user', "|")

# Step 3. Assign it to a variable called users
# setting the index of the Data Frame to "user_id"
users = df1.set_index(['user_id'])
users.head()
```

```
Out[161]:
```

	age	gender	occupation	zip_code
user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213

```
In [162]: # Step 4. Discover what is the mean age per occupation

users.groupby('occupation').mean(['age'])
```

```
Out[162]:
```

	age
occupation	
administrator	38.746835
artist	31.392857
doctor	43.571429
educator	42.010526
engineer	36.388060
entertainment	29.222222
executive	38.718750
healthcare	41.562500
homemaker	32.571429
lawyer	36.750000
librarian	40.000000
marketing	37.615385
none	26.555556
other	34.523810
programmer	33.121212
retired	63.071429
salesman	35.666667
scientist	35.548387

age  
occupation

student 22021622

```
In [166... # Step 5. Discover the Male ratio per occupation and sort it from the most to the least  
maleRatio = pd.pivot_table(users, aggfunc = 'count', index = 'occupation', values = 'age', column
```

```
In [168... # Step 5 (cont)  
  
# determine the total number of individuals of each occupation  
total = maleRatio[['M', 'F']].sum(axis = 1)  
  
maleRatio['maleRatio'] = (maleRatio['M'] / total)  
  
maleRatio.sort_values(by = ['maleRatio'], ascending = False)
```

```
Out[168... gender  F  M  maleRatio  
occupation  
doctor      0   7  1.000000  
engineer    2  65  0.970149  
technician  1  26  0.962963  
retired     1  13  0.928571  
programmer  6  60  0.909091  
executive   3  29  0.906250  
scientist   3  28  0.903226  
entertainment 2  16  0.888889  
lawyer      2  10  0.833333  
salesman    3   9  0.750000  
educator    26  69  0.726316  
student     60 136  0.693878  
other       36  69  0.657143  
marketing   10  16  0.615385  
writer      19  26  0.577778  
none        4   5  0.555556  
administrator 36  43  0.544304  
artist      13  15  0.535714  
librarian   29  22  0.431373  
healthcare  11   5  0.312500  
homemaker   6   1  0.142857
```

```
In [169... # Step 6. For each occupation, calculate the minimum and maximum ages  
users.groupby('occupation').agg({'age': ['min', 'max']})
```

```
Out[169... age  
min  max  
occupation  
administrator  21  70  
artist         19  48  
doctor         28  64  
educator       23  63
```

	age	
	min	max
<b>occupation</b>		
engineer	22	70
entertainment	15	50
executive	22	69
healthcare	22	62
homemaker	20	50
lawyer	21	53
librarian	23	69
marketing	24	55
none	11	55
other	13	64
programmer	20	63
retired	51	73
salesman	18	66
scientist	23	55

```
In [170... # Step 7. For each combination of occupation and sex, calculate the mean age

users.groupby(['occupation', 'gender']).agg({'age': 'mean'})
```

Out[170...

occupation	gender	age
administrator	F	40.638889
	M	37.162791
artist	F	30.307692
	M	32.333333
doctor	M	43.571429
educator	F	39.115385
	M	43.101449
engineer	F	29.500000
	M	36.600000
entertainment	F	31.000000
	M	29.000000
executive	F	44.000000
	M	38.172414
healthcare	F	39.818182
	M	45.400000
homemaker	F	34.166667
	M	23.000000
lawyer	F	39.500000
	M	36.200000
librarian	F	40.000000
	M	40.000000
marketing	F	37.200000
	M	37.875000

		age
occupation		gender
none	F	36.500000
	M	18.600000
other	F	35.472222
	M	34.028986
programmer	F	32.166667
	M	33.216667
retired	F	70.000000
	M	62.538462
salesman	F	27.000000
	M	38.555556
scientist	F	28.333333
	M	36.321429
student	F	20.750000
	M	22.669118
technician	F	38.000000
	M	32.961538

In [171]...

```
# Step 8. For each occupation present the percentage of women and men

# references:
# https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.round.html
# https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.drop.html

# reusing some of the code from step 5:
maleRatio = pd.pivot_table(users, aggfunc = 'count', index = 'occupation', values = 'age', column
# determine the percentage of males -> multiple ratio by 100 to get %
maleRatio['male %'] = (maleRatio['M'] / total) * 100

# determine percentage of females
maleRatio['male %'] = (1 - (maleRatio['M'] / total)) * 100

# round percentage to the nearest tenth
# remove raw values from the table
maleRatio.round(decimals = 1).drop(columns = ['M', 'F'], axis = 0)
```

Out[171]...

gender	male %
occupation	
administrator	45.6
artist	46.4
doctor	0.0
educator	27.4
engineer	3.0
entertainment	11.1
executive	9.4
healthcare	68.8
homemaker	85.7
lawyer	16.7
librarian	56.9
marketing	38.5
none	44.4
other	34.3

gender	male %
occupation	
programmer	9.1
retired	7.1
salesman	25.0
scientist	9.7
student	30.6

## Question 2 Euro Teams - Python with Pandas

```
In [173... # Step 1. Import the necessary libraries
import pandas as pd
import numpy as np

# Step 2. Import the dataset from this address
# Step 3. Assign it to a variable called euro12
euro12 = pd.read_csv(r'https://raw.githubusercontent.com/guipsamora/pandas_exercises/master/02_FilesandIO/euro12.csv')

euro12.head()
```

```
Out[173...
   Team  Goals  Shots  Shots  Shooting  %  Total  Hit  Penalty  Penalties  Saves  Saves-  Fouls  Fouls
   Team  Goals  on  off  Accuracy  Goals-  shots  Woodwork  goals  not  made  to-  Won  Conc
   Team  Goals  target  target  %  to-  (inc.  %  scored  ...  ratio  %  %
   Team  Goals  target  target  %  to-  (inc.  %  scored  ...  ratio  %  %
0  Croatia    4    13    12    51.9%  16.0%    32     0     0     0  ...  13  81.3%  41
1  Czech    4    13    18    41.9%  12.9%    39     0     0     0  ...  9  60.1%  53
   Republic
2  Denmark    4    10    10    50.0%  20.0%    27     1     0     0  ...  10  66.7%  25
3  England    5    11    18    50.0%  17.2%    40     0     0     0  ...  22  88.1%  43
4  France     3    22    24    37.9%   6.5%    65     1     0     0  ...   6  54.6%  36
```

5 rows × 35 columns

```
In [174... # Step 4. Select only the Goal column

euro12['Goals']
```

```
Out[174...
0    4
1    4
2    4
3    5
4    3
5   10
6    5
7    6
8    2
9    2
10   6
11   1
12   5
13  12
14   5
15   2
Name: Goals, dtype: int64
```

```
In [175... # Step 5. How many team participated in the Euro2012?

euro12['Team'].count()
```

```
Out[175... 16
```

```
In [176... # Step 6. What is the number of columns in the dataset?

len(euro12.columns)
```

Out[176... 35

```
In [177... # Step 7. View only the columns Team, Yellow Cards and Red Cards and assign them to a dataframe c

discipline = euro12[['Team', 'Yellow Cards', 'Red Cards']]
discipline
```

Out[177...

	Team	Yellow Cards	Red Cards
0	Croatia	9	0
1	Czech Republic	7	0
2	Denmark	4	0
3	England	5	0
4	France	6	0
5	Germany	4	0
6	Greece	9	1
7	Italy	16	0
8	Netherlands	5	0
9	Poland	7	1
10	Portugal	12	0
11	Republic of Ireland	6	1
12	Russia	6	0
13	Spain	11	0
14	Sweden	7	0
15	Ukraine	5	0

```
In [178... # Step 8. Sort the teams by Red Cards, then to Yellow Cards

discipline.sort_values(['Red Cards', 'Yellow Cards'], ascending = [True, True])
```

Out[178...

	Team	Yellow Cards	Red Cards
2	Denmark	4	0
5	Germany	4	0
3	England	5	0
8	Netherlands	5	0
15	Ukraine	5	0
4	France	6	0
12	Russia	6	0
1	Czech Republic	7	0
14	Sweden	7	0
0	Croatia	9	0
13	Spain	11	0
10	Portugal	12	0
7	Italy	16	0
11	Republic of Ireland	6	1
9	Poland	7	1
6	Greece	9	1

```
In [179... # Step 9. Calculate the mean Yellow Cards given per Team

euro12.groupby('Team').agg({'Yellow Cards': 'mean'})
```

Out[179... **Yellow Cards**

Team	
Croatia	9
Czech Republic	7
Denmark	4
England	5
France	6
Germany	4
Greece	9
Italy	16
Netherlands	5
Poland	7
Portugal	12
Republic of Ireland	6
Russia	6
Spain	11
Sweden	7
Ukraine	5

```
In [180... # Step 10. Filter teams that scored more than 6 goals

euro12.loc[euro12['Goals'] > 6]
```

Out[180...

	Team	Goals	Shots on target	Shots off target	Shooting Accuracy	Goals-to-shots	% Total shots (inc. Blocked)	Hit Woodwork	Penalty goals	Penalties not scored	...	Saves made	Saves-to-shots ratio	Fouls Won	Conc
5	Germany	10	32	32	47.8%	15.6%	80	2	1	0	...	10	62.6%	63	
13	Spain	12	42	33	55.9%	16.0%	100	0	1	0	...	15	93.8%	102	

2 rows × 35 columns

```
In [181... # goalsStep 11. Select the teams that start with G

euro12[euro12['Team'].str.startswith('G')]
```

Out[181...

	Team	Goals	Shots on target	Shots off target	Shooting Accuracy	Goals-to-shots	% Total shots (inc. Blocked)	Hit Woodwork	Penalty goals	Penalties not scored	...	Saves made	Saves-to-shots ratio	Fouls Won	Conc
5	Germany	10	32	32	47.8%	15.6%	80	2	1	0	...	10	62.6%	63	
6	Greece	5	8	18	30.7%	19.2%	32	1	1	1	...	13	65.1%	67	

2 rows × 35 columns

```
In [182... # Step 12. Select the first 7 columns

euro12[euro12.columns[0:7]]
```

Out[182...

	Team	Goals	Shots on	Shots off	Shooting	% Goals-to-	Total shots (inc.
--	------	-------	----------	-----------	----------	-------------	-------------------

			target	target	Accuracy	shots	Blocked)
0	Croatia	4	13	12	51.9%	16.0%	32
1	Czech Republic	4	13	18	41.9%	12.9%	39
2	Denmark	4	10	10	50.0%	20.0%	27
3	England	5	11	18	50.0%	17.2%	40
4	France	3	22	24	37.9%	6.5%	65
5	Germany	10	32	32	47.8%	15.6%	80
6	Greece	5	8	18	30.7%	19.2%	32
7	Italy	6	34	45	43.0%	7.5%	110
8	Netherlands	2	12	36	25.0%	4.1%	60
9	Poland	2	15	23	39.4%	5.2%	48
10	Portugal	6	22	42	34.3%	9.3%	82
11	Republic of Ireland	1	7	12	36.8%	5.2%	28
12	Russia	5	9	31	22.5%	12.5%	59
13	Spain	12	42	33	55.9%	16.0%	100
14	Sweden	5	17	19	47.2%	13.8%	39
15	Ukraine	2	7	26	21.2%	6.0%	38

```
In [183... # Step 13. Select all columns except the last 3
euro12[euro12.columns[:-3]]
```

	Team	Goals	Shots on target	Shots off target	Shooting Accuracy	% Goals-to-shots	Total shots (inc. Blocked)	Hit Woodwork	Penalty goals	Penalties not scored	...	Clean Sheets	Blocks	Goals conceded
0	Croatia	4	13	12	51.9%	16.0%	32	0	0	0	...	0	10	
1	Czech Republic	4	13	18	41.9%	12.9%	39	0	0	0	...	1	10	
2	Denmark	4	10	10	50.0%	20.0%	27	1	0	0	...	1	10	
3	England	5	11	18	50.0%	17.2%	40	0	0	0	...	2	29	
4	France	3	22	24	37.9%	6.5%	65	1	0	0	...	1	7	
5	Germany	10	32	32	47.8%	15.6%	80	2	1	0	...	1	11	
6	Greece	5	8	18	30.7%	19.2%	32	1	1	1	...	1	23	
7	Italy	6	34	45	43.0%	7.5%	110	2	0	0	...	2	18	
8	Netherlands	2	12	36	25.0%	4.1%	60	2	0	0	...	0	9	
9	Poland	2	15	23	39.4%	5.2%	48	0	0	0	...	0	8	
10	Portugal	6	22	42	34.3%	9.3%	82	6	0	0	...	2	11	
11	Republic of Ireland	1	7	12	36.8%	5.2%	28	0	0	0	...	0	23	
12	Russia	5	9	31	22.5%	12.5%	59	2	0	0	...	0	8	
13	Spain	12	42	33	55.9%	16.0%	100	0	1	0	...	5	8	
14	Sweden	5	17	19	47.2%	13.8%	39	3	0	0	...	1	12	
15	Ukraine	2	7	26	21.2%	6.0%	38	0	0	0	...	0	4	

16 rows × 32 columns

```
In [184... # Step 14. Present only the Shooting Accuracy from England, Italy and Russia
euro12[euro12.Team.isin(['England','Italy','Russia'])][['Team','Shooting Accuracy']]
```



Out[184...

	Team	Shooting Accuracy
3	England	50.0%
7	Italy	43.0%
12	Russia	22.5%

## Question 3 Housing - Python with Pandas

In [185...

```
# Step 1. Import the necessary libraries
import numpy as np
import pandas as pd
import random

# Step 2. Create 3 different Series, each of length 100, as follows:
# • The first a random number from 1 to 4

series1 = pd.Series(np.random.randint(1,5, size = 100))
series1
```

Out[185...

```
0      2
1      1
2      1
3      2
4      2
..
95     3
96     3
97     4
98     3
99     2
Length: 100, dtype: int32
```

In [186...

```
# Step 2. Create 3 different Series, each of length 100, as follows:
# • The second a random number from 1 to 3

series2 = pd.Series(np.random.randint(1,4, size = 100))
series2
```

Out[186...

```
0      3
1      3
2      3
3      2
4      2
..
95     3
96     2
97     1
98     1
99     3
Length: 100, dtype: int32
```

In [187...

```
# Step 2. Create 3 different Series, each of length 100, as follows:
# • The third a random number from 10,000 to 30,000

series3 = pd.Series(np.random.randint(10000,30001, size = 100))
series3
```

Out[187...

```
0      19542
1      28145
2      12789
3      18538
4      25725
...
95      11041
96      22085
97      29262
98      12849
99      28913
Length: 100, dtype: int32
```

```
In [188... # Step 3. Create a DataFrame by joining the Series by column

housing_data_frame = pd.DataFrame({'series1': series1, 'series2': series2, 'series3': series3})
housing_data_frame
```

```
Out[188...      series1  series2  series3
0         2         3   19542
1         1         3   28145
2         1         3   12789
3         2         2   18538
4         2         2   25725
...      ...      ...      ...
95        3         3   11041
96        3         2   22085
97        4         1   29262
98        3         1   12849
99        2         3   28913
```

100 rows × 3 columns

```
In [189... # Step 4. Change the name of the columns to bedrs, bathrs, price_sqr_meter

# add all three series together in a dictionary
# assign a column label to each series

housing_data_frame.columns = ["bedrs", "bathrs", "price_sqr_meter"]
housing_data_frame
```

```
Out[189...      bedrs  bathrs  price_sqr_meter
0         2         3          19542
1         1         3          28145
2         1         3          12789
3         2         2          18538
4         2         2          25725
...      ...      ...              ...
95        3         3          11041
96        3         2          22085
97        4         1          29262
98        3         1          12849
99        2         3          28913
```

100 rows × 3 columns

```
In [190... # Step 5. Create a one column DataFrame with the values of the 3 Series and assign it to 'bigcolumn'

bigcolumn = pd.concat([series1, series2, series3])
bigcolumn
```

```
Out[190... 0         2
1         1
2         1
3         2
4         2
...
95      11041
96      22085
97      29262
```

```
98     12849
99     28913
100     28913
```

```
In [191... # Step 6. Ops it seems it is going only until index 99. Is it true?

# due to the concatenation from the previous step, the indices for each series is only considered
# showing unique indices. That's why instead of showing an index of 299, it's showing us an index

# We can use these two codes to double check this:
print("The length is:", len(bigcolumn))

if (max(bigcolumn.index)==99):
    print("True, the index is 99")
else:
    print("False, the index is not 99")
```

```
The length is: 300
True, the index is 99
```

```
In [192... # Step 7. Reindex the DataFrame so it goes from 0 to 299

# Reference:
# https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.reset_index.html?highlight=reset_

bigcolumn.reset_index(drop = True)
```

```
Out[192... 0         2
1         1
2         1
3         2
4         2
...
295     11041
296     22085
297     29262
298     12849
299     28913
Length: 300, dtype: int32
```

## Question 4 Wind Statistics - Python with Pandas

```
In [193... # Step 1. Import the necessary libraries
import pandas as pd
import numpy as np
import datetime

# Step 2. Import the dataset from this address
wind_stats_data = pd.read_csv(r'https://raw.githubusercontent.com/guipsamora/pandas_exercises/master/
wind_stats_data = wind_stats_data.rename(columns = {'Yr': 'Year', 'Mo': 'Month', 'Dy': 'Day'})

wind_stats_data
```

```
Out[193...
```

	Year	Month	Day	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAL
0	61	1	1	15.04	14.96	13.17	9.29	NaN	9.87	13.67	10.25	10.83	12.58	18.50	15.04
1	61	1	2	14.71	NaN	10.83	6.50	12.62	7.67	11.50	10.04	9.79	9.67	17.54	13.83
2	61	1	3	18.50	16.88	12.33	10.13	11.17	6.17	11.25	NaN	8.50	7.67	12.75	12.71
3	61	1	4	10.58	6.63	11.75	4.58	4.54	2.88	8.63	1.79	5.83	5.88	5.46	10.88
4	61	1	5	13.33	13.25	11.42	6.17	10.71	8.21	11.92	6.54	10.92	10.34	12.92	11.83
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
6569	78	12	27	17.58	16.96	17.62	8.08	13.21	11.67	14.46	15.59	14.04	14.00	17.21	40.08
6570	78	12	28	13.21	5.46	13.46	5.00	8.12	9.42	14.33	16.25	15.25	18.05	21.79	41.46
6571	78	12	29	14.00	10.29	14.42	8.71	9.71	10.54	19.17	12.46	14.50	16.42	18.88	29.58
6572	78	12	30	18.50	14.04	21.29	9.13	12.75	9.71	18.08	12.87	12.46	12.12	14.67	28.79
6573	78	12	31	20.33	17.41	27.29	9.59	12.08	10.13	19.25	11.63	11.58	11.38	12.08	22.08

6574 rows × 15 columns

```
In [194... # Step 3. Assign it to a variable called data and replace the first 3 columns by a proper datetime

wind_stats_data["Date"] = pd.to_datetime(wind_stats_data[['Year', 'Month', 'Day']].astype(str).agg
wind_stats_data = wind_stats_data.drop(columns = ['Year', 'Month', 'Day'])
column_names = ["Date", "RPT", "VAL", "ROS", "KIL", "SHA", "BIR", "DUB", "CLA", "MUL", "CLO", "BEL
wind_stats_data = wind_stats_data.reindex(columns = column_names)
wind_stats_data
```

```
Out[194...      Date  RPT  VAL  ROS  KIL  SHA  BIR  DUB  CLA  MUL  CLO  BEL  MAL
0  2061-01-01  15.04  14.96  13.17  9.29  NaN  9.87  13.67  10.25  10.83  12.58  18.50  15.04
1  2061-01-02  14.71   NaN  10.83  6.50  12.62  7.67  11.50  10.04  9.79  9.67  17.54  13.83
2  2061-01-03  18.50  16.88  12.33  10.13  11.17  6.17  11.25   NaN  8.50  7.67  12.75  12.71
3  2061-01-04  10.58  6.63  11.75  4.58  4.54  2.88  8.63  1.79  5.83  5.88  5.46  10.88
4  2061-01-05  13.33  13.25  11.42  6.17  10.71  8.21  11.92  6.54  10.92  10.34  12.92  11.83
...      ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...
6569  1978-12-27  17.58  16.96  17.62  8.08  13.21  11.67  14.46  15.59  14.04  14.00  17.21  40.08
6570  1978-12-28  13.21  5.46  13.46  5.00  8.12  9.42  14.33  16.25  15.25  18.05  21.79  41.46
6571  1978-12-29  14.00  10.29  14.42  8.71  9.71  10.54  19.17  12.46  14.50  16.42  18.88  29.58
6572  1978-12-30  18.50  14.04  21.29  9.13  12.75  9.71  18.08  12.87  12.46  12.12  14.67  28.79
6573  1978-12-31  20.33  17.41  27.29  9.59  12.08  10.13  19.25  11.63  11.58  11.38  12.08  22.08
```

6574 rows × 13 columns

```
In [196... # Step 4. Year 2061? Do we really have data from this year? Create a function to fix it and apply

def correct_date(col_name):

    if col_name.year > 2000:
        year = col_name.year - 100
    else:
        year = col_name.year
    return datetime.date(year, col_name.month, col_name.day)
```

```
In [197... # step 4 (cont) (DONE)

wind_stats_data["Date"] = wind_stats_data["Date"].apply(correct_date)
wind_stats_data
```

```
Out[197...      Date  RPT  VAL  ROS  KIL  SHA  BIR  DUB  CLA  MUL  CLO  BEL  MAL
0  1961-01-01  15.04  14.96  13.17  9.29  NaN  9.87  13.67  10.25  10.83  12.58  18.50  15.04
1  1961-01-02  14.71   NaN  10.83  6.50  12.62  7.67  11.50  10.04  9.79  9.67  17.54  13.83
2  1961-01-03  18.50  16.88  12.33  10.13  11.17  6.17  11.25   NaN  8.50  7.67  12.75  12.71
3  1961-01-04  10.58  6.63  11.75  4.58  4.54  2.88  8.63  1.79  5.83  5.88  5.46  10.88
4  1961-01-05  13.33  13.25  11.42  6.17  10.71  8.21  11.92  6.54  10.92  10.34  12.92  11.83
...      ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...
6569  1978-12-27  17.58  16.96  17.62  8.08  13.21  11.67  14.46  15.59  14.04  14.00  17.21  40.08
6570  1978-12-28  13.21  5.46  13.46  5.00  8.12  9.42  14.33  16.25  15.25  18.05  21.79  41.46
6571  1978-12-29  14.00  10.29  14.42  8.71  9.71  10.54  19.17  12.46  14.50  16.42  18.88  29.58
6572  1978-12-30  18.50  14.04  21.29  9.13  12.75  9.71  18.08  12.87  12.46  12.12  14.67  28.79
6573  1978-12-31  20.33  17.41  27.29  9.59  12.08  10.13  19.25  11.63  11.58  11.38  12.08  22.08
```

6574 rows × 13 columns

```
In [198... # Step 5. Set the right dates as the index. Pay attention at the data type, it should be datetime

wind_stats_data_new = wind_stats_data.set_index("Date")
wind_stats_data_new.index.astype("datetime64[ns]")
```

```
Out[198... DatetimeIndex(['1961-01-01', '1961-01-02', '1961-01-03', '1961-01-04',
                  '1961-01-05', '1961-01-06', '1961-01-07', '1961-01-08',
                  '1961-01-09', '1961-01-10',
                  ...
                  '1978-12-22', '1978-12-23', '1978-12-24', '1978-12-25',
                  '1978-12-26', '1978-12-27', '1978-12-28', '1978-12-29',
                  '1978-12-30', '1978-12-31'],
                  dtype='datetime64[ns]', name='Date', length=6574, freq=None)
```

```
In [199... # Step 6. Compute how many values are missing for each location over the entire record. They should
# all calculations below.

wind_stats_data_new.isnull().sum()
```

```
Out[199... RPT      6
          VAL      3
          ROS      2
          KIL      5
          SHA      2
          BIR      0
          DUB      3
          CLA      2
          MUL      3
          CLO      1
          BEL      0
          MAL      4
          dtype: int64
```

```
In [200... # Step 7. Compute how many non-missing values there are in total.

wind_stats_data_new.count()
```

```
Out[200... RPT      6568
          VAL      6571
          ROS      6572
          KIL      6569
          SHA      6572
          BIR      6574
          DUB      6571
          CLA      6572
          MUL      6571
          CLO      6573
          BEL      6574
          MAL      6570
          dtype: int64
```

```
In [201... # Step 8. Calculate the mean windspeeds of the windspeeds over all the locations and all the time

wind_stats_data_new.mean()
```

```
Out[201... RPT      12.362987
          VAL      10.644314
          ROS      11.660526
          KIL       6.306468
          SHA      10.455834
          BIR       7.092254
          DUB       9.797343
          CLA       8.495053
          MUL       8.493590
          CLO       8.707332
          BEL      13.121007
          MAL      15.599079
          dtype: float64
```

In [202...

```
# Step 9. Create a DataFrame called loc_stats and calculate the min, max and mean windspeeds and
# of the windspeeds at each location over all the days

loc_stats = pd.DataFrame()
loc_stats['min'] = wind_stats_data_new.min(axis = 0)
loc_stats['max'] = wind_stats_data_new.max(axis = 0)
loc_stats['mean'] = wind_stats_data_new.mean(axis = 0)
loc_stats['std'] = wind_stats_data_new.std(axis = 0)
loc_stats
```

Out[202...

	min	max	mean	std
<b>RPT</b>	0.67	35.80	12.362987	5.618413
<b>VAL</b>	0.21	33.37	10.644314	5.267356
<b>ROS</b>	1.50	33.84	11.660526	5.008450
<b>KIL</b>	0.00	28.46	6.306468	3.605811
<b>SHA</b>	0.13	37.54	10.455834	4.936125
<b>BIR</b>	0.00	26.16	7.092254	3.968683
<b>DUB</b>	0.00	30.37	9.797343	4.977555
<b>CLA</b>	0.00	31.08	8.495053	4.499449
<b>MUL</b>	0.00	25.88	8.493590	4.166872
<b>CLO</b>	0.04	28.21	8.707332	4.503954
<b>BEL</b>	0.13	42.38	13.121007	5.835037
<b>MAL</b>	0.67	42.54	15.599079	6.699794

In [203...

```
# Step 10. Create a DataFrame called day_stats and calculate the min, max and mean windspeed and
# of the windspeeds across all the locations at each day.

day_stats = pd.DataFrame()
day_stats['min'] = wind_stats_data_new.min(axis = 1)
day_stats['max'] = wind_stats_data_new.max(axis = 1)
day_stats['mean'] = wind_stats_data_new.mean(axis = 1)
day_stats['std'] = wind_stats_data_new.std(axis = 1)
day_stats
```

Out[203...

	min	max	mean	std
<b>Date</b>				
<b>1961-01-01</b>	9.29	18.50	13.018182	2.808875
<b>1961-01-02</b>	6.50	17.54	11.336364	3.188994
<b>1961-01-03</b>	6.17	18.50	11.641818	3.681912
<b>1961-01-04</b>	1.79	11.75	6.619167	3.198126
<b>1961-01-05</b>	6.17	13.33	10.630000	2.445356
...	...	...	...	...
<b>1978-12-27</b>	8.08	40.08	16.708333	7.868076
<b>1978-12-28</b>	5.00	41.46	15.150000	9.687857
<b>1978-12-29</b>	8.71	29.58	14.890000	5.756836
<b>1978-12-30</b>	9.13	28.79	15.367500	5.540437
<b>1978-12-31</b>	9.59	27.29	15.402500	5.702483

6574 rows × 4 columns

In [204...

```
# Step 11. Find the average windspeed in January for each location.
# Treat January 1961 and January 1962 both as January.

wind_stats_data['month'] = pd.DatetimeIndex(wind_stats_data['Date']).month
january_avg = wind_stats_data.where(wind_stats_data['month'] == 1)
january_avg.loc[:, 'RPT': 'MAL'].mean()
```

Out[204...

```
RPT    14.847325
VAL    12.914560
ROS    13.299624
KIL     7.199498
SHA    11.667734
BIR     8.054839
DUB    11.819355
CLA     9.512047
MUL     9.543208
CLO    10.053566
BEL    14.550520
MAL    18.028763
dtype: float64
```

In [205...

```
# Step 12. Downsample the record to a yearly frequency for each location.

wind_stats_data_new.asfreq('Y')
```

Out[205...

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAL
Date												
1961-12-31	9.87	7.83	7.67	3.75	5.66	3.50	10.04	3.08	5.04	3.79	8.04	14.67
1962-12-31	22.67	16.88	28.67	14.12	19.75	17.08	27.79	25.21	19.83	17.79	25.46	37.63
1963-12-31	13.88	14.42	12.12	9.25	14.33	10.67	18.29	11.96	12.04	15.37	16.79	14.09
1964-12-31	16.33	19.25	13.37	10.08	17.04	12.54	19.83	13.79	12.67	15.04	21.37	23.58
1965-12-31	13.62	13.88	12.29	6.08	12.33	7.41	9.59	10.21	7.46	12.17	15.71	16.75
1966-12-31	13.00	11.46	10.13	6.34	11.87	7.50	13.50	8.46	11.00	10.04	17.29	22.46
1967-12-31	16.88	13.75	11.34	9.08	13.54	7.71	11.75	11.83	11.83	11.75	17.25	22.63
1968-12-31	9.13	2.13	7.38	2.50	4.04	0.50	6.83	2.54	3.54	5.50	5.71	12.42
1969-12-31	14.42	13.83	27.71	7.08	12.08	10.00	14.58	11.00	12.54	7.12	11.17	17.41
1970-12-31	8.38	0.37	9.59	2.62	1.75	0.08	4.83	2.13	2.54	1.17	3.67	7.21
1971-12-31	14.88	10.50	26.08	8.46	13.50	10.04	21.04	10.25	13.54	11.34	12.12	27.33
1972-12-31	13.83	14.46	15.87	9.75	8.71	11.00	10.67	11.54	11.50	10.75	18.00	17.50
1973-12-31	10.67	10.04	6.87	1.46	6.96	5.75	3.83	6.21	4.75	6.13	12.79	15.79
1974-12-31	16.04	16.29	15.21	8.42	13.67	9.75	15.25	16.13	15.04	13.46	18.54	18.46
1975-12-31	15.59	12.33	13.42	2.37	4.08	1.17	7.08	4.25	5.91	6.34	11.38	19.55
1976-12-31	8.67	8.83	9.38	3.67	5.37	4.58	7.92	1.79	4.46	4.38	6.38	15.67
1977-12-31	15.09	7.62	8.79	7.08	10.63	7.58	15.59	11.54	12.25	9.08	14.12	19.55
1978-12-31	20.33	17.41	27.29	9.59	12.08	10.13	19.25	11.63	11.58	11.38	12.08	22.08

In [206...

```
# Step 13. Downsample the record to a monthly frequency for each location.

wind_stats_data_new.asfreq('M')
```

Out[206...

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAL
Date												
1961-01-31	24.21	19.55	16.71	11.96	14.42	10.46	14.88	8.21	10.50	9.96	12.42	13.92
1961-02-28	12.92	12.75	NaN	8.92	16.13	12.29	14.75	14.46	13.96	14.04	18.41	13.17
1961-03-31	8.96	8.04	9.13	8.50	10.75	9.54	11.92	9.59	11.25	8.54	11.96	12.21
1961-04-30	11.67	11.00	9.54	5.54	9.42	5.79	5.09	8.25	6.96	6.25	12.21	8.75
1961-05-31	7.00	9.79	12.25	4.83	8.25	5.37	6.58	9.29	6.58	7.12	11.87	10.63

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAL
Date												
...	...	...	...	...	...	...	...	...	...	...	...	...
1978-08-31	11.54	5.54	7.41	4.67	7.62	6.17	8.87	5.25	7.83	6.17	11.58	16.88
1978-09-30	26.75	15.63	16.54	13.37	17.58	13.13	16.92	13.79	13.46	13.79	18.91	31.88
1978-10-31	8.58	4.29	10.79	4.29	4.08	2.71	4.63	1.04	3.67	2.75	8.71	10.67
1978-11-30	15.34	4.54	14.75	3.50	4.54	4.96	7.50	2.42	4.96	3.75	4.92	11.50
1978-12-31	20.33	17.41	27.29	9.59	12.08	10.13	19.25	11.63	11.58	11.38	12.08	22.08

```
In [207... # Step 14. Downsample the record to a weekly frequency for each location.

wind_stats_data_new. asfreq('W')
```

	RPT	VAL	ROS	KIL	SHA	BIR	DUB	CLA	MUL	CLO	BEL	MAL
Date												
1961-01-01	15.04	14.96	13.17	9.29	NaN	9.87	13.67	10.25	10.83	12.58	18.50	15.04
1961-01-08	10.96	9.75	7.62	5.91	9.62	7.29	14.29	7.62	9.25	10.46	16.62	16.46
1961-01-15	12.04	9.67	11.75	2.37	7.38	3.13	2.50	6.83	4.75	5.63	7.54	6.75
1961-01-22	9.59	5.88	9.92	2.17	6.87	5.50	9.38	7.04	6.34	7.50	10.88	9.92
1961-01-29	NaN	23.91	22.29	17.54	24.08	19.70	22.00	20.25	21.46	19.95	27.71	23.38
...	...	...	...	...	...	...	...	...	...	...	...	...
1978-12-03	21.21	21.34	17.75	11.58	16.75	14.46	17.46	15.29	15.79	17.50	21.42	25.75
1978-12-10	24.92	22.54	16.54	14.62	15.59	13.00	13.21	14.12	16.21	16.17	26.08	21.92
1978-12-17	9.87	3.21	8.04	2.21	3.04	0.54	2.46	1.46	1.29	2.67	5.00	9.08
1978-12-24	8.67	5.63	12.12	4.79	5.09	5.91	12.25	9.25	10.83	11.71	11.92	31.71
1978-12-31	20.33	17.41	27.29	9.59	12.08	10.13	19.25	11.63	11.58	11.38	12.08	22.08

940 rows × 12 columns

```
In [208... # Step 15. Calculate the min, max and mean windspeeds and standard deviations of the windspeeds a
# for each week (assume that the first week starts on January 2 1961) for the first 52 weeks.

df = wind_stats_data_new[wind_stats_data_new.index < pd.to_datetime('1962-01-01')]
df. asfreq('W'). mean()
df. asfreq('W'). min()
df. asfreq('W'). max()
df. asfreq('W'). std()

day_stats
```

	min	max	mean	std
Date				
1961-01-01	9.29	18.50	13.018182	2.808875
1961-01-02	6.50	17.54	11.336364	3.188994
1961-01-03	6.17	18.50	11.641818	3.681912
1961-01-04	1.79	11.75	6.619167	3.198126
1961-01-05	6.17	13.33	10.630000	2.445356
...	...	...	...	...
1978-12-27	8.08	40.08	16.708333	7.868076
1978-12-28	5.00	41.46	15.150000	9.687857
1978-12-29	8.71	29.58	14.890000	5.756836
1978-12-30	9.13	28.79	15.367500	5.540437



min max mean std  
Date

## Question 5 Food - Python with Pandas

```
In [225... # Step 1. Import the necessary libraries
import pandas as pd
import numpy as np

# Step 2. Import the dataset from this address.
# Step 3. Assign it to a variable called chipo.
chipo = pd.read_csv(r'https://raw.githubusercontent.com/justmarkham/DAT8/master/data/chipotle.tsv')

# Step 4. See the first 10 entries
chipo.head(10)
```

```
Out[225... order_id  quantity  item_name  choice_description  item_price
0          1          1  Chips and Fresh Tomato Salsa          NaN          $2.39
1          1          1                lzze  [Clementine]          $3.39
2          1          1  Nantucket Nectar  [Apple]          $3.39
3          1          1  Chips and Tomatillo-Green Chili Salsa          NaN          $2.39
4          2          2  Chicken Bowl  [Tomatillo-Red Chili Salsa (Hot), [Black Beans...          $16.98
5          3          1  Chicken Bowl  [Fresh Tomato Salsa (Mild), [Rice, Cheese, Sou...          $10.98
6          3          1  Side of Chips          NaN          $1.69
7          4          1  Steak Burrito  [Tomatillo Red Chili Salsa, [Fajita Vegetables...          $11.75
8          4          1  Steak Soft Tacos  [Tomatillo Green Chili Salsa, [Pinto Beans, Ch...          $9.25
9          5          1  Steak Burrito  [Fresh Tomato Salsa, [Rice, Black Beans, Pinto...          $9.25
```

```
In [226... # Step 5. What is the number of observations in the dataset?

len(chipo)
```

Out[226... 4622

```
In [227... # Step 6. What is the number of columns in the dataset?

len(chipo.columns)
```

Out[227... 5

```
In [228... # Step 7. Print the name of all the columns.

chipo.columns
```

```
Out[228... Index(['order_id', 'quantity', 'item_name', 'choice_description',
      'item_price'],
      dtype='object')
```

```
In [229... # Step 8. How is the dataset indexed?

chipo.index
```

Out[229... RangeIndex(start=0, stop=4622, step=1)

```
In [230... # Step 9. Which was the most-ordered item?

most_ordered_item = chipo['item_name'].value_counts()
most_ordered_item

# as we can see, "Chicken Bowl" is the most ordered item with 726 orders.
```

```
Out[230... Chicken Bowl      726
Chicken Burrito    553
Chips and Guacamole 479
Steak Burrito      368
Canned Soft Drink  301
Chips              211
Steak Bowl         211
Bottled Water      162
Chicken Soft Tacos  115
Chips and Fresh Tomato Salsa 110
Chicken Salad Bowl 110
Canned Soda        104
Side of Chips      101
Veggie Burrito     95
Barbacoa Burrito   91
Veggie Bowl        85
Carnitas Bowl      68
Barbacoa Bowl      66
Carnitas Burrito   59
Steak Soft Tacos   55
6 Pack Soft Drink  54
Chips and Tomatillo Red Chili Salsa 48
Chicken Crispy Tacos 47
Chips and Tomatillo Green Chili Salsa 43
Carnitas Soft Tacos 40
Steak Crispy Tacos 35
Chips and Tomatillo-Green Chili Salsa 31
Steak Salad Bowl   29
Nantucket Nectar   27
Barbacoa Soft Tacos 25
Chips and Roasted Chili Corn Salsa 22
Izze               20
Chips and Tomatillo-Red Chili Salsa 20
Chips and Roasted Chili-Corn Salsa 18
Veggie Salad Bowl  18
Barbacoa Crispy Tacos 11
Barbacoa Salad Bowl 10
Chicken Salad       9
Veggie Soft Tacos   7
Carnitas Crispy Tacos 7
Carnitas Salad Bowl 6
Veggie Salad        6
Burrito            6
Steak Salad         4
Salad              2
Bowl               2
Crispy Tacos        2
Veggie Crispy Tacos 1
Chips and Mild Fresh Tomato Salsa 1
Carnitas Salad      1
Name: item_name, dtype: int64
```

```
In [231... # Step 10. For the most-ordered item, how many items were ordered?

most_ordered_item[:1]
```

```
Out[231... Chicken Bowl      726
Name: item_name, dtype: int64
```

```
In [232... # Step 11. What was the most ordered item in the choice_description column?

chipo.choice_description.value_counts()

# as we can see, Diet Code is the most ordered item with 134 orders.
```

```
Out[232... [Diet Coke]
134
[Coke]
123
[Sprite]
77
[Fresh Tomato Salsa, [Rice, Black Beans, Cheese, Sour Cream, Lettuce]]
42
[Fresh Tomato Salsa, [Rice, Black Beans, Cheese, Sour Cream, Guacamole, Lettuce]]
```

40

```
...
[Fresh Tomato (Mild), [Rice, Sour Cream, Cheese]]
1
[Tomatillo Red Chili Salsa, [Rice, Cheese, Guacamole]]
1
[Tomatillo Red Chili Salsa, [Rice, Fajita Vegetables, Black Beans, Cheese, Lettuce]]
1
[Fresh Tomato Salsa (Mild), [Fajita Veggies, Cheese, Sour Cream, Guacamole, Lettuce]]
1
[[Roasted Chili Corn Salsa (Medium), Tomatillo-Red Chili Salsa (Hot)], [Black Beans, Rice, Fajita
Veggies, Cheese, Sour Cream, Guacamole, Lettuce]] 1
Name: order_id, dtype: int64
```

```
In [233... # Step 12. How many items were orderd in total?

chipo['quantity'].sum()
```

Out[233... 4972

```
In [234... # Step 13.
# • Turn the item price into a float
# • Check the item price type
# • Create a lambda function and change the type of item price
# • Check the item price type

# turn the item into a float
chipo['item_price'] = chipo['item_price'].apply(lambda x: float(x[1:]))
chipo['item_price'].dtypes
```

Out[234... dtype('float64')

```
In [235... # Step 14. How much was the revenue for the period in the dataset?

chipo['revenue'] = chipo['quantity'] * chipo['item_price']
chipo['revenue'].sum()
```

Out[235... 39237.02

```
In [236... # Step 15. How many orders were made in the period?

orders = chipo['order_id'].nunique()
orders
```

Out[236... 1834

```
In [237... # Step 16. What is the average revenue amount per order?

chipo.groupby('order_id')['revenue'].mean()
```

```
Out[237... order_id
1      2.890000
2     33.960000
3      6.335000
4     10.500000
5      6.850000
...
1830   11.500000
1831    4.300000
1832    6.600000
1833   11.750000
1834    9.583333
Name: revenue, Length: 1834, dtype: float64
```

```
In [238... # Step 17. How many different items are sold?

chipo['item_name'].nunique()
```

Out[238... 50

## Question 6

Create a line plot showing the number of marriages and divorces per capita in the U.S. between 1867 and 2014. Label both lines and show the legend.

Don't forget to label your axes!

```
In [1]: # import pandas to first inspect the data
import pandas as pd

marriages_divorces = pd.read_csv('us-marriages-divorces-1867-2014.csv')
marriages_divorces.head()
```

```
Out[1]:
```

	Year	Marriages	Divorces	Population	Marriages_per_1000	Divorces_per_1000
0	1867	357000.0	10000.0	36970000	9.7	0.3
1	1868	345000.0	10000.0	37885000	9.1	0.3
2	1869	348000.0	11000.0	38870000	9.0	0.3
3	1870	352000.0	11000.0	39905000	8.8	0.3
4	1871	359000.0	12000.0	41010000	8.8	0.3

```
In [2]: # import other libraries
import matplotlib.pyplot as plt
%matplotlib inline

marriages_divorces.plot.line(x = 'Year', y = ['Marriages_per_1000', 'Divorces_per_1000'], figsize

# label the plot and axes


# x-axis
plt.xlabel('Year', fontsize = 14)

# y-axis
plt.ylabel('Marriages & Divorces Per Capita', fontsize = 14)

# add the title
plt.title('# of U.S Marriages and Divorces per Capita: 1867 - 2014', fontsize = 18)

# show the legend
plt.legend(prop = dict(size = 14))

plt.show()
```



## Question 7

Create a vertical bar chart comparing the number of marriages and divorces per capita in the U.S. between 1900, 1950,

and 2000.

Don't forget to label your axes!

```
In [3]: # only the years we need
marriages_divorces = marriages_divorces[(marriages_divorces.Year == 1900) | (marriages_divorces.Y
```

```
In [4]: # import other libraries
import matplotlib.pyplot as plt
%matplotlib inline

marriages_divorces = marriages_divorces.drop(columns = ['Marriages', 'Divorces', 'Population'])
marriages_divorces
marriages_divorces = marriages_divorces.set_index('Year')

marriages_divorces.plot.bar(figsize = (10,6), color = {"#2a9d8f", "#e9c46a"})

# label the chart and axes


# x-axis
plt.xlabel('Year', fontsize = 14)

# y-axis
plt.ylabel('Marriages & Divorces', fontsize = 18)

# add the title
plt.title('# of U.S Marriages and Divorces per Capita: 1900, 1950, 2000')

# show the legend
plt.legend(prop = dict(size = 14))

plt.show()
```



## Question 8

Create a horizontal bar chart that compares the deadliest actors in Hollywood. Sort the actors by their kill count and label each bar with the corresponding actor's name.

Don't forget to label your axes!

```
In [5]: # import to create the visual
import matplotlib.pyplot as plt
%matplotlib inline

# import pandas to use the Data Frame
import pandas as pd

# import the data into the Data Frame
hollywood_kills = pd.read_csv('actor_kill_counts.csv')
hollywood_kills = hollywood_kills.sort_values('Count')

# add the title
plt.title("Hollywood's Top 10 Deadliest Actors", fontsize = 18)

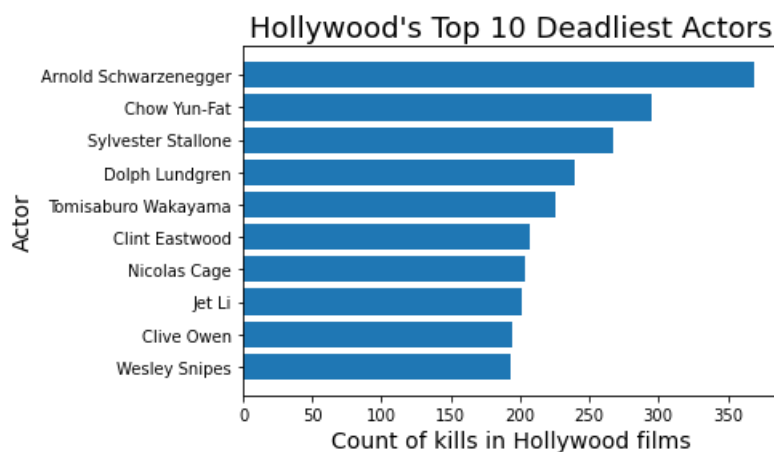
# label the chart and axes

# x-axis
plt.ylabel('Actor', fontsize = 14)

# y-axis
plt.xlabel('Count of kills in Hollywood films', fontsize = 14)

plt.barh(hollywood_kills['Actor'], hollywood_kills.Count)

plt.show()
```



## Question 9

Create a pie chart showing the fraction of all Roman Emperors that were assassinated.

Make sure that the pie chart is an even circle, labels the categories, and shows the percentage breakdown of the categories.

```
In [6]: # import pandas to use the Data Frame
import pandas as pd

# add data to the Data Frame
roman_emp = pd.read_csv('roman-emperor-reigns.csv')
roman_emp.head()
```

```
Out[6]:
```

	Emperor	Length_of_Reign	Cause_of_Death
0	Augustus	40.58	Possibly assassinated
1	Tiberius	22.50	Possibly assassinated
2	Caligula	4.83	Assassinated
3	Claudius	13.75	Possibly assassinated
4	Nero	13.67	Suicide

```
In [7]: # group by cause of death and sum totals of each death type
roman_emp_death = roman_emp.groupby('Cause_of_Death').count().drop(columns = 'Length_of_Reign')

# sort them in the decreasing order
roman_emp_death.sort_values(by = ['Emperor'], ascending = False)
```

```
Out[7]:
```

Cause_of_Death	Emperor
Assassinated	22
Natural causes	16
Killed in battle	8
Possibly assassinated	8
Illness	5
Suicide	5
Executed	3
Died in captivity	1

```
In [8]: # import other libraries
import matplotlib.pyplot as plt
import numpy as np


# specify the data
y = np.array([22, 16, 8, 8, 5, 5, 3, 1])
label = ["Assassinated", 'Natural causes', 'Killed in battle', 'Possibly assassinated',
'Illness', 'Suicide', 'Executed', 'Died in captivity']

# explode on "Assassinated"
explode = (0.05, 0, 0, 0, 0, 0, 0, 0)

# create the plot
fig1, ax1 = plt.subplots(figsize = (8,8))

# set the labels, axes, and title
ax1.axis('equal')
ax1.pie(y, labels = label, autopct = '%1.1f%%', explode = explode)
ax1.set_title('Roman Emperors Cause of Death', fontsize = 18)

plt.show()
```



## Question 10

Create a scatter plot showing the relationship between the total revenue earned by arcades and the number of Computer Science PhDs awarded in the U.S. between 2000 and 2009.

Don't forget to label your axes!

Color each dot according to its year.

```
In [9]: # import pandas to use the Data Frame
import pandas as pd

# import the csv file and add it to the Data Frame
arcade_science_phd = pd.read_csv('arcade-revenue-vs-cs-doctorates.csv')
arcade_science_phd
```

```
Out[9]:
```

	Year	Total Arcade Revenue (billions)	Computer Science Doctorates Awarded (US)
0	2000	1.196	861
1	2001	1.176	830
2	2002	1.269	809
3	2003	1.240	867
4	2004	1.307	948
5	2005	1.435	1129
6	2006	1.601	1453
7	2007	1.654	1656
8	2008	1.803	1787
9	2009	1.734	1611

```
In [10]: # import other libraries
import seaborn as sns
import matplotlib.pyplot as plt

# increase the size of the figure
plt.figure(figsize = (10,6))

# create the scatter plot, set the title, and increase the size of the data points
sns.scatterplot(x = 'Computer Science Doctorates Awarded (US)',
                y = 'Total Arcade Revenue (billions)',
                data = arcade_science_phd, hue = 'Year', s = 150,).set_title('Total Annual Comput

# label the axes
plt.xlabel('Year', fontsize = 14)
plt.ylabel('Total Arcade Revenue ($ billions)', fontsize = 14)
```

```
Out[10]: Text(0, 0.5, 'Total Arcade Revenue ($ billions)')
```



## Total Annual Computer Science PhD Awarded vs. Total Arcade Revenues



In [ ]: